

Graphical Analysis in ArtToMusic

Rafael De Smet

December 06, 2016

Contents

1	Introduction	2
2	Algorithms	2
2.1	Edge Detection	2
2.1.1	Convolution kernel	2
2.1.2	Sobel	4
2.1.3	Frei-Chen	5
2.1.4	Prewitt	6
2.1.5	Roberts Cross	7
2.1.6	LoG	7
2.1.7	Scharr	8
2.2	Color Analysis	9
2.2.1	RGB - HSV - CMYK	9
2.2.2	Black And White Balance	10
2.2.3	Color Quantization - Image Segmentation	11
2.3	Image Hashing	11
2.3.1	Average Hashing	11
2.3.2	Difference Hash	12
2.3.3	Perceptive Hash	12
2.3.4	Use Of The Hash	12
2.4	Entropy	12

1 Introduction

In this paper I will be discussing several graphical analysis algorithms. Some of these are actually used in the code of ArtToMusic and others are discussed out of interest. I will also compare the algorithms I used and didn't use. These algorithms will return data which is used to generate music.

2 Algorithms

2.1 Edge Detection

The concept of edge detection algorithms is pretty straightforward. They try to detect all the edges in an image, edges are the contours of objects, shapes in the image. Edge detection algorithms all use what are called convolution kernels. A kernel in image processing is a small matrix used to apply effects to an image, such as blurring, outlining. Here we will see kernels used for edge detection only. Listed below are six of the best and most used algorithms.

- Sobel
- Frei-Chen
- Prewitt
- Roberts Cross
- LoG
- Scharr

2.1.1 Convolution kernel

Since all the algorithms are based on the mathematical principle of convolution, an explanation of these convolution kernels is in order.

Convolution is the technique of multiplying together two arrays of different size but of the same dimensionality. An array of dimensionality two is simply a matrix. When working with images the pixels are represented as a (2D) matrix and the kernel is also a 2D matrix. The kernel is a small matrix that we will multiply with the image matrix to perform the convolution. This

kernel matrix is different for each edge detection algorithm and we will see examples of different matrices in later chapters of this paper.

The idea of these algorithms is to create a new image that shows us the edges of the image. Each pixel of the original image is added to its local neighbours, weighted by the kernel. This produces a new image. If the kernel is chosen wisely, we get all the edges found in the image. Mathematically we can write the convolution as follows, with O the output image, I the input image and K the kernel. $I(i, j)$ means the pixel on the i^{th} row and the j^{th} column.

$$O(i, j) = \sum_{k=1}^m \sum_{l=1}^n I(i + k - 1, j + l - 1) K(k, l) \quad (1)$$

An example will clarify the previous.

I₁₁	I₁₂	I₁₃	I₁₄	I₁₅	I₁₆	I₁₇	I₁₈	I₁₉
I₂₁	I₂₂	I₂₃	I₂₄	I₂₅	I₂₆	I₂₇	I₂₈	I₂₉
I₃₁	I₃₂	I₃₃	I₃₄	I₃₅	I₃₆	I₃₇	I₃₈	I₃₉
I₄₁	I₄₂	I₄₃	I₄₄	I₄₅	I₄₆	I₄₇	I₄₈	I₄₉
I₅₁	I₅₂	I₅₃	I₅₄	I₅₅	I₅₆	I₅₇	I₅₈	I₅₉
I₆₁	I₆₂	I₆₃	I₆₄	I₆₅	I₆₆	I₆₇	I₆₈	I₆₉

K₁₁	K₁₂	K₁₃
K₂₁	K₂₂	K₂₃

Figure 1: A pixel matrix and a kernel matrix

We expect an output image that shows us the edges of the original image. Using the kernel matrix we will compute every new pixel of the output image, by sliding the kernel matrix over the original pixels. Each kernel position corresponds to a single output pixel, the value of which is calculated by equation (1).

In our example, the value of the bottom right pixel in the output image will be found as follows [2]:

$$O_{57} = I_{57}K_{11} + I_{58}K_{12} + I_{59}K_{13} + I_{67}K_{21} + I_{68}K_{22} + I_{69}K_{23} \quad (2)$$

In figure 2 you can see an idyllic landscape. We will use this image to see how well the different edge detection algorithms work.



Figure 2: Idyllic landscape

2.1.2 Sobel

The Sobel algorithm performs a 2D spatial gradient measurement and finds regions of 'high spatial frequency' or edges. It uses two 3x3 kernels, one for the vertical edges and the other for the horizontal edges. These two kernels can be applied separately and afterwards combined to find all the edges. In figure 3 you can see the result of the Sobel filter.

The horizontal kernel: $\begin{vmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{vmatrix}$ and the vertical kernel: $\begin{vmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{vmatrix}$



Figure 3: Sobel filter applied to Figure 2

2.1.3 Frei-Chen

The Frei-Chen algorithm also uses 3x3 kernels, but this time there are nine different convolution kernels. The four first matrices, G_1 , G_2 , G_3 , G_4 , are used for edges, the next four are used for lines and the last one is used to compute averages. Figure 4 shows the result of the Frei-Chen filter.

$$\begin{aligned}
 G_1 &= \frac{1}{2\sqrt{2}} \begin{vmatrix} 1 & \sqrt{2} & 1 \\ 0 & 0 & 0 \\ -1 & -\sqrt{2} & -1 \end{vmatrix} & G_2 &= \frac{1}{2\sqrt{2}} \begin{vmatrix} 1 & 0 & -1 \\ \sqrt{2} & 0 & -\sqrt{2} \\ 1 & 0 & -1 \end{vmatrix} & G_3 &= \frac{1}{2\sqrt{2}} \begin{vmatrix} 0 & -1 & \sqrt{2} \\ 1 & 0 & -1 \\ -\sqrt{2} & 1 & 0 \end{vmatrix} \\
 G_4 &= \frac{1}{2\sqrt{2}} \begin{vmatrix} \sqrt{2} & -1 & 0 \\ -1 & 0 & 1 \\ 0 & 1 & -\sqrt{2} \end{vmatrix} & G_5 &= \frac{1}{2} \begin{vmatrix} 0 & 1 & 0 \\ -1 & 0 & -1 \\ 0 & 1 & 0 \end{vmatrix} & G_6 &= \frac{1}{2} \begin{vmatrix} -1 & 0 & 1 \\ 0 & 0 & 0 \\ 1 & 0 & -1 \end{vmatrix} \\
 G_7 &= \frac{1}{6} \begin{vmatrix} 1 & -2 & 1 \\ -2 & 4 & -2 \\ 1 & -2 & 1 \end{vmatrix} & G_8 &= \frac{1}{6} \begin{vmatrix} -2 & 1 & -2 \\ 1 & 4 & 1 \\ -2 & 1 & -2 \end{vmatrix} & G_9 &= \frac{1}{3} \begin{vmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{vmatrix}
 \end{aligned}$$



Figure 4: Frei-Chen filter applied to Figure 2

2.1.4 Prewitt

This algorithm is very similar to the Sobel algorithm. Again, two kernels are used, one for the horizontal and one for the vertical edges. Afterwards, they are recombined to get all the edges in the image. But in this case the kernels are considerably simpler. In figure 5 you can see the result of this filter.

$$\text{Horizontal filter} = \begin{vmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{vmatrix} \quad \text{Vertical filter} = \begin{vmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{vmatrix}$$



Figure 5: Prewitt filter applied to Figure 2

2.1.5 Roberts Cross

This algorithm uses even simpler kernels than Prewitt does. This time we use two 2x2 kernels. These kernels correspond to the edges running at 45° to the pixel grid, one for each of the two perpendicular orientations. Figure 6 shows the result of the Roberts Cross filter.

$$\text{Horizontal filter} = \begin{vmatrix} 1 & 0 \\ 0 & -1 \end{vmatrix} \quad \text{Vertical filter} = \begin{vmatrix} 0 & 1 \\ -1 & 0 \end{vmatrix}$$



Figure 6: Roberts Cross filter applied to Figure 2

2.1.6 LoG

This algorithm combines two methods, the Gaussian filtering method¹ and the Laplacian method for edge detection². Hence the name "Laplacian of Gaussian" (LoG). The edge points of an image are detected by finding the zero crossings of the 2nd derivative of the image intensity. Because the 2nd derivative is very sensitive to noise, which could give us bad results, the Gaussian filter is used to clear the noise from the image. In figure 7, you can find the result of this filter.

¹The Gaussian filter is used to blur images and remove noise and detail.

²This method highlights regions in the image of rapid intensity change, so it is useful for edge detection.

The R library OpenImageR³ uses the following LoG mask.

$$\text{LoG mask} = \begin{vmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{vmatrix}$$




Figure 7: LoG filter applied to Figure 2

2.1.7 Scharr

This algorithm is an extension of the Sobel algorithm. Although the Sobel kernels are very good, they do not have perfect rotational symmetry. This is what the Scharr kernels try to optimize. Figure 8 shows the results.

The most frequently used kernels are the following:

$$\text{Horizontal filter} = \begin{vmatrix} 3 & 10 & 3 \\ 0 & 0 & 0 \\ -3 & -10 & -3 \end{vmatrix} \quad \text{Vertical filter} = \begin{vmatrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{vmatrix}$$

³<https://cran.r-project.org/web/packages/OpenImageR/OpenImageR.pdf>



Figure 8: Scharr filter applied to Figure 2

2.2 Color Analysis

Only edge detection algorithms are not sufficient to get enough data from the image to form a musical interpretation. We will need some form of color analysis. This section describes some algorithms and methods we can use to obtain information about the colors of the image.

2.2.1 RGB - HSV - CMYK

To represent color on the computer the RGB representation is the most commonly used. Each pixel is described using three values, the amount of red (R), green (G) and blue (B). Using these values we can count how many pixels in the image are dominantly red, green or blue. When we're going to translate the image data into musical patterns, we can match each color to another kind of music (happy, sad, etc...)

HSV, sometimes called HSB, is another representation of a pixel. This time we use the hue (H), which color it is, the saturation (S) and the brightness or value (B or V) to describe the pixel. Where the RGB model consists of three values indicating the amounts of a certain color each pixel has, the HSV model consists of three independent components to form a color.

Hue First we have the hue, which is basically the color. Here the color is represented using a circle with all the colors on it. The value of this component is the degrees of the angle we have to make on the circle to get this color.

Saturation The saturation indicates the fullness of the color. This value is expressed in a percentage, with 0% a gray, flat color and 100% a full, rich color.

Value of Brightness The value or brightness indicates the lightness of the color and is also expressed in a percentage, 0% is black and 100% is white.

This representation will be used in combination with the RGB model in the translation into the musical patterns. In figure 9 you can see how the hue, saturation and value are combined and calculated.




Figure 9: The HSV color representation

CMYK is another representation of color, based on the mixing of four colors, cyan (C), magenta (M), yellow (Y) and key (K, black). This model is used frequently with printing, but not very interesting for this paper.

2.2.2 Black And White Balance

During the image analysis we will convert the image to a gray scale image, so we can determine the amount of white and black in it. Using the RGB model, this is an easy process. Black in the RGB model is (0,0,0) and white is (255,255,255). So we can simply count the amount of pixels that are very close to those two values and we know the amount of black and white pixels

in the image. Analogous to this we can count every gray pixel in the image, by finding every pixel where the RGB values are exactly the same.

2.2.3 Color Quantization - Image Segmentation

Sometimes it can be easier to work with images that are partitioned in simpler regions. This technique is called image segmentation or color quantification. There are many different ways to implement this technique. One such simple approach is using the k-means cluster algorithm. This algorithm attempts to partition a data set into k clusters. The data set contains the RGB values of each pixel. The disadvantage of this algorithm is that we don't know what the right value for k is. So we have to perform this algorithm a couple of times, each time with a different k to see which value gives us a good result.

A second technique is the X-means algorithm which is an improvement of the k-means algorithm. Where we don't know how many clusters there are when using k-means, and thus guessing, the X-means algorithm can determine the number of clusters by itself. Based on the total number of clusters found by X-means, we can generate parts of the piece of music.

2.3 Image Hashing

It is also possible to get the hash value of an image. The hash is a hexadecimal number.

2.3.1 Average Hashing

The first hash method is the average hash (aHash) of the image. This algorithm works in five steps.

1. Convert the image to grayscale.
2. Reduce the size of the image, to reduce the number of computations.
3. Average the resulting colors. For an 8x8 image, 64 will be averaged.
4. Compute the bits of the hash value by comparing if each color value is above or below the mean.
5. Construct the hash.

2.3.2 Difference Hash

We can also compute the dHash of an image. It is similar to the average hash, but now the difference between adjacent pixels is also considered in the computation.

2.3.3 Perceptive Hash

This method uses the discrete cosine transform (DCT) and compares the pixels based on the frequencies, given by the DCT, instead of the color values. DCT is a similar technique as Fourier analysis, it expresses a finite sequence of points, in this case the pixels, in terms of a sum of cosine functions.

2.3.4 Use Of The Hash

One application of these methods is the recognition of images. Search engines like Google use this technique to search similar images to any image you provide to the search engine. I could use this as a way to store information about the image my application has already scanned and generated music for.

For example, if you use a picture of a tree with the ArtToMusic application, it will produce a certain kind of music. If you use another picture of another tree that looks quite similar to the first tree, you expect to get a similar result in the music. The hash values of the images will help with this.

2.4 Entropy

Another method we can use to get information from an image is to calculate the entropy. The entropy of an image is a measure of the amount of disorder in the image. This technique is also used in other disciplines, for example to study the structures of living organisms.

When talking about images, we consider the following. If all the pixels of the image have the same level, the entropy is zero. This means there is not a lot of information (almost none) to be gained from the image. An image of a blue sky has a very low entropy cause we only see blue.

When all the pixels in the image are different, the entropy of the image is maximum and there is a lot of information to get from the image. As the entropy describes how much (dis)order there is, we will use this to determine how much (dis)order there is in the generated music.

To get the right information about the pixels, we use a histogram that shows the count of the distinct pixel values. The entropy of an image H is defined as:

$$H = - \sum_{k=0}^{M-1} p_k \log_2(p_k) \quad (3)$$

where M is the number of unique pixel values and p_k is the count of each pixel level, how many times this pixel level occurs in the image.

You can calculate the entropy of an image, but it is more useful to calculate the entropy of an analysis of that picture. This gives us information about the analysis and not just the image in general.

References

- [1] Lampros Mouselimis *OpenPackageR, An Image Processing Toolkit* 2017
- [2] R. Fisher, S. Perkins, A. Walker, E. Wolfart *Hypermedia Image Processing Reference* 2003
- [3] Daniel Rákos' blog (<http://rastergrid.com/blog/2011/01/frei-chen-edge-detector/>) 2011
- [4] University of Auckland, New Zealand, Computer Science course Computer Graphics and Image Processing *Gaussian Filtering* 2010
- [5] Computação Visual e Multimédia, course on Image Processing
- [6] RoboRealm blog <http://www.roborealm.com/help/Prewitt.php> 2017
- [7] Simon Colton, Pedro Torres *Evolving Approximate Image Filters* 2009, Computational Creativity Group Department of Computing, Imperial College London

- [8] Dr. Neal Krawetz' blog (<http://www.hackerfactor.com/blog/?/archives/529-Kind-of-Like-That.html>) 2013
- [9] John Loomis' site (<http://www.johnloomis.org/ece563/notes/basics/entropy/entropy.html>) 1998
- [10] Kevin Meurer's blog (<http://kevinmeurer.com/a-simple-guide-to-entropy-based-discretization/>) 2015
- [11] Ryan Walker's blogpost (<https://www.r-bloggers.com/color-quantization-in-r/>) 2016